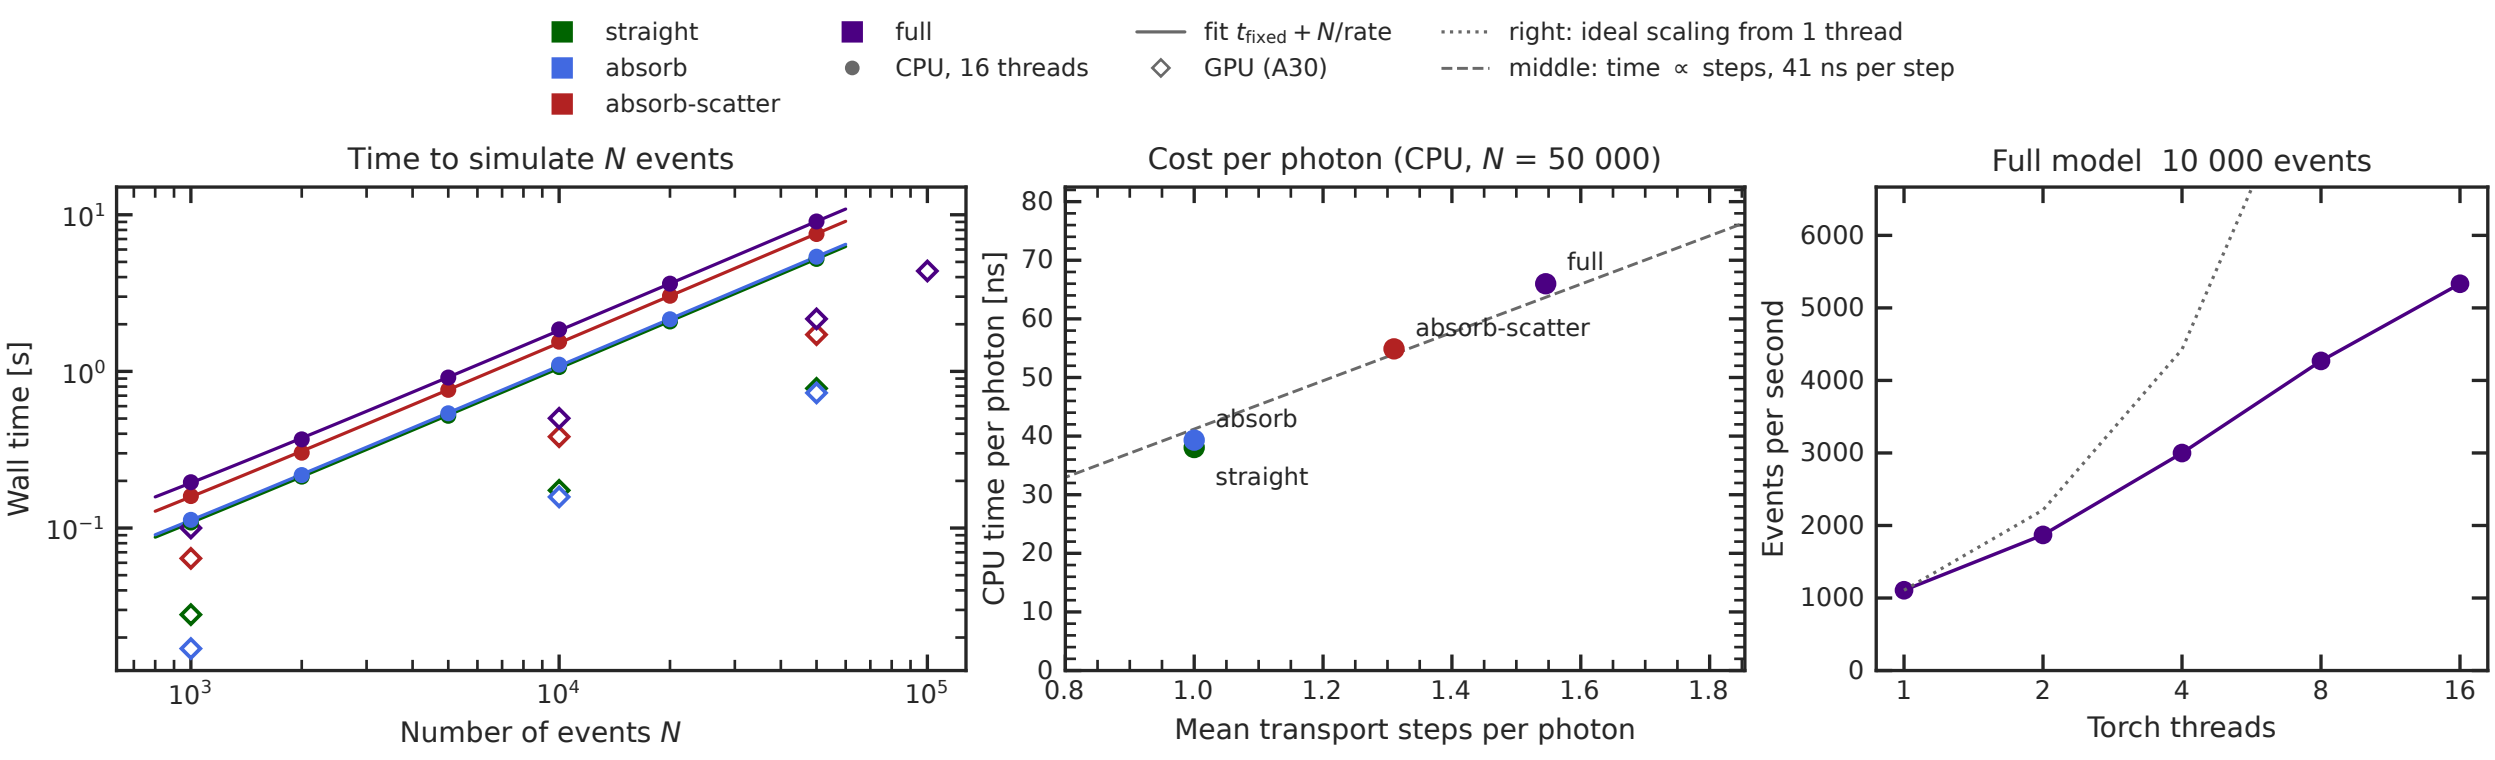
straight
absorb
absorb-scatter
full
CPU, 16 threads
fit $t_{fixed} + N/rate$
GPU (A30)
right: ideal scaling from 1 thread
middle: time ∝ steps, 41 ns per step
Time to simulate N events
Wall time [s]
Number of events N
Cost per photon (CPU, N = 50 000)
CPU time per photon [ns]
Mean transport steps per photon
full
absorb-scatter
absorb
straight
Full model 10 000 events
Events per second
Torch threads